

# final-S1G1(1)

May 17, 2020

## 1 Final Exam: Part B

- Section: Sec01
- Group Number: S1G1: Mercedes Bischoff, Ryan Mandel, Bryce Owen, Kimberly Tang
- Due Date: May 7, 2020
- Purpose: Culmination of Semester Knowledge Base

### 1.1 Preprocessing data

```
[80]: ##--Load Libraries  
import pandas as pd  
import seaborn as sns  
import numpy as np  
import matplotlib.pyplot as plt  
from scipy import stats  
from sklearn.model_selection import train_test_split  
from sklearn.neural_network import MLPClassifier  
from sklearn import metrics
```

```
[81]: ##--Import and Preview Data Set  
plants = pd.read_excel('http://barney.gonzaga.edu/~chuang/data/plants.xlsx')  
plants.head()
```

```
[81]:   ID  sepal_length  sepal_width  petal_length  petal_width  species  
0    1           4.3           2.8           1.9           NaN         0  
1    2           4.8           3.4           1.6           0.2         0  
2    3           5.0           3.0           1.8           0.4         0  
3    4           4.8           3.0           1.4           0.3         0  
4    5           5.1           3.8           1.6           0.2         0
```

```
[82]: ##--Identify Missing Data  
plants.isnull().sum()
```

```
[82]: ID                0  
sepal_length         1  
sepal_width          5  
petal_length        10
```

```

petal_width      8
species          0
dtype: int64

```

```

[83]: #--Identify records that have missing values
plants[plants.isnull().any(axis=1)]

```

```

[83]:      ID  sepal_length  sepal_width  petal_length  petal_width  species
0      1           4.3           2.8           1.9           NaN           0
32     33           5.4           2.9           1.7           NaN           0
35     36           5.7           3.6           NaN           0.2           0
38     39           5.3           NaN           1.4           0.2           0
39     40           4.7           NaN           NaN           NaN           0
46     47           4.6           3.0           NaN           0.5           0
62     63           NaN           3.1           1.8           0.3           0
67     68           5.2           3.8           NaN           0.4           0
102   103           5.2           2.7           NaN           0.5           0
115   116           5.5           2.4           3.7           NaN           1
142   143           5.1           NaN           4.4           1.6           1
166   167           5.3           3.0           NaN           1.8           1
191   192           4.9           NaN           NaN           1.0           1
199   200           5.8           2.7           4.1           NaN           1
228   229           5.8           2.4           NaN           2.1           2
239   240           5.8           NaN           5.1           1.9           2
262   263           5.9           3.6           NaN           1.5           2
266   267           6.4           3.2           6.3           NaN           2
297   298           5.7           2.5           5.0           NaN           2
316   317           7.3           2.9           NaN           1.8           2
323   324           7.6           2.4           6.8           NaN           2

```

```

[84]: #Fill all missing data with averages grouped by species
plants['sepal_length'] = plants['sepal_length'].fillna(plants.
    ↳groupby('species')['sepal_length'].transform('mean'))
plants['sepal_width'] = plants['sepal_width'].fillna(plants.
    ↳groupby('species')['sepal_width'].transform('mean'))
plants['petal_length'] = plants['petal_length'].fillna(plants.
    ↳groupby('species')['petal_length'].transform('mean'))
plants['petal_width'] = plants['petal_width'].fillna(plants.
    ↳groupby('species')['petal_width'].transform('mean'))
plants.head()

```

```

[84]:      ID  sepal_length  sepal_width  petal_length  petal_width  species
0      1           4.3           2.8           1.9      0.299057           0
1      2           4.8           3.4           1.6      0.200000           0
2      3           5.0           3.0           1.8      0.400000           0
3      4           4.8           3.0           1.4      0.300000           0
4      5           5.1           3.8           1.6      0.200000           0

```

```
[85]: ##-- Standardize numeric variable by adding zscore columns
plants['sepal_length_z'] = stats.zscore(plants['sepal_length'], nan_policy =_
↳ 'omit')
plants['sepal_width_z'] = stats.zscore(plants['sepal_width'], nan_policy =_
↳ 'omit')
plants['petal_length_z'] = stats.zscore(plants['petal_length'], nan_policy =_
↳ 'omit')
plants['petal_width_z'] = stats.zscore(plants['petal_width'], nan_policy =_
↳ 'omit')

plants.head()
```

```
[85]:   ID  sepal_length  sepal_width  petal_length  petal_width  species  \
0    1           4.3           2.8           1.9    0.299057         0
1    2           4.8           3.4           1.6    0.200000         0
2    3           5.0           3.0           1.8    0.400000         0
3    4           4.8           3.0           1.4    0.300000         0
4    5           5.1           3.8           1.6    0.200000         0

      sepal_length_z  sepal_width_z  petal_length_z  petal_width_z
0          -1.679790         -0.436789         -1.034591         -1.234960
1          -1.121177          0.703984         -1.201071         -1.365184
2          -0.897731         -0.056531         -1.090085         -1.102255
3          -1.121177         -0.056531         -1.312058         -1.233719
4          -0.786009          1.464498         -1.201071         -1.365184
```

```
[86]: ##--Identify rows with zscores >3 or <-3
plants[(plants['sepal_length_z'] > 3) |
      (plants['sepal_length_z'] < -3) |
      (plants['sepal_width_z'] > 3) |
      (plants['sepal_width_z'] < -3) |
      (plants['petal_length_z'] > 3) |
      (plants['petal_length_z'] < -3) |
      (plants['petal_width_z'] > 3) |
      (plants['petal_width_z'] < -3)]
```

```
[86]:   ID  sepal_length  sepal_width  petal_length  petal_width  species  \
106  107      4.800000      5.020648           1.8    0.200000         0
116  117      8.900000      2.600000           3.3    1.300000         1
154  155      8.975003      2.900000           4.9    1.500000         1
216  217      7.000000      4.974328           5.1    1.800000         1
261  262      9.157097      3.300000           5.2    2.400000         2
282  283      8.975003      3.800000           6.7    2.200000         2
294  295      6.700000      2.700000           5.0    4.052461         2

      sepal_length_z  sepal_width_z  petal_length_z  petal_width_z
106          -1.121177          3.785301         -1.090085         -1.365184
116           3.459455         -0.817046         -0.257685          0.080925
```

154	3.543250	-0.246660	0.630209	0.343854
216	1.336723	3.697233	0.741195	0.738247
261	3.746691	0.513855	0.796688	1.527034
282	3.543250	1.464498	1.629088	1.264105
294	1.001555	-0.626917	0.685702	3.699432

```
[87]: #--Drop rows that contain outliers
plants = plants.drop(plants[(plants['sepal_length_z'] > 3) |
    (plants['sepal_length_z'] < -3) |
    (plants['sepal_width_z'] > 3) |
    (plants['sepal_width_z'] < -3) |
    (plants['petal_length_z'] > 3) |
    (plants['petal_length_z'] < -3) |
    (plants['petal_width_z'] > 3) |
    (plants['petal_width_z'] < -3)].index)
plants.shape
```

```
[87]: (319, 10)
```

```
[88]: #Find duplicate records
plants[plants.duplicated(subset=plants.columns.difference(['ID']))]
```

```
[88]:
```

	ID	sepal_length	sepal_width	petal_length	petal_width	species	\
54	55	4.9	3.1	1.5	0.1	0	
91	92	4.9	3.1	1.5	0.1	0	

	sepal_length_z	sepal_width_z	petal_length_z	petal_width_z
54	-1.009454	0.133597	-1.256565	-1.496648
91	-1.009454	0.133597	-1.256565	-1.496648

```
[89]: #Drop those records
plants.drop_duplicates(subset=plants.columns.difference(['ID']), inplace = True)
plants.shape
```

```
[89]: (317, 10)
```

```
[90]: #--Add Species Name Column to Data Set
def Species(x):
    if x == 0:
        return "setosa"
    elif x == 1:
        return "versicolor"
    else:
        return "virginica"

plants['SpeciesName'] = plants['species'].apply(Species)
plants.head()
```

```
[90]:
```






	ID	sepal_length	sepal_width	petal_length	petal_width	species	\
0	1	4.3	2.8	1.9	0.299057	0	
1	2	4.8	3.4	1.6	0.200000	0	

2	3	5.0	3.0	1.8	0.400000	0
3	4	4.8	3.0	1.4	0.300000	0
4	5	5.1	3.8	1.6	0.200000	0

	sepal_length_z	sepal_width_z	petal_length_z	petal_width_z	SpeciesName
0	-1.679790	-0.436789	-1.034591	-1.234960	setosa
1	-1.121177	0.703984	-1.201071	-1.365184	setosa
2	-0.897731	-0.056531	-1.090085	-1.102255	setosa
3	-1.121177	-0.056531	-1.312058	-1.233719	setosa
4	-0.786009	1.464498	-1.201071	-1.365184	setosa

## 1.2 Explore Dataset

```
[91]: #--Histogram of Each Variable (Z score variables not needed because they are
      ↪the same distribution as normal variables)
plants[['petal_length','petal_width','sepal_length','sepal_width','species']].
      ↪hist(bins=10,figsize=(20,15))
plt.show()
```



```
[92]: #--Correlation Chart (Z score variables not needed because they are the same  
      →distribution as normal variables)  
plants[['petal_length', 'petal_width', 'sepal_length', 'sepal_width', 'species']].  
      →corr().style.background_gradient("Greens")
```

```
[92]: <pandas.io.formats.style.Styler at 0x12b50af60>
```


### 1.3 Clustering

```
[93]: #Choose variabes to cluster and justify choice  
      # We chose to select the two variables with the strongest correlation to  
      →species, petal_length and petal_width  
  
X = plants[['petal_length', 'petal_width']]  
X.head()
```

```
[93]:   petal_length  petal_width  
0          1.9      0.299057  
1          1.6      0.200000  
2          1.8      0.400000  
3          1.4      0.300000  
4          1.6      0.200000
```

```
[94]: #Plot the distribution of sepal_length and petal_length  
plt.figure(figsize=(8,6))  
sns.scatterplot(X['petal_length'], X['petal_width'])
```

```
[94]: <matplotlib.axes._subplots.AxesSubplot at 0x12b632048>
```



```
[95]: #Save means and standard deviations
      #Used to standardize data of new plants for prediction
plants_mean = X.mean()
plants_std = X.std()
print(plants_mean)
print()
print(plants_std)
```

```
petal_length    3.760817
petal_width     1.230537
dtype: float64
```

```
petal_length    1.804528
petal_width     0.742804
dtype: float64
```

```
[96]: #Standardize the data
z_score = stats.zscore(X)
X_z = pd.DataFrame(z_score, columns = ['petal_length_z', 'petal_width_z'])
X_z.tail()
```

[96]:	petal_length_z	petal_width_z
312	1.686861	1.576880
313	0.687793	0.363339
314	1.686861	1.085716
315	0.965312	0.767852
316	0.743297	1.576880

```
[97]: #Fit a model with the data
      #Create three clusters
      from sklearn.cluster import KMeans

      kmeans_plants = KMeans(n_clusters = 3).fit(X_z)
```

```
[98]: #Obtain the labels
cluster = kmeans_plants.labels_
cluster
```

```
[98]: array([1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,  
            1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,  
            1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,  
            1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,  
            1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,  
            1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 2, 2, 2, 2,  
            2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2,  
            2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2,  
            2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2,  
            2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2,  
            2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2,  
            2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 0, 0, 0, 0, 0, 2, 0, 0,  
            0, 2, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 2, 0, 0, 0, 0, 0, 0,  
            2, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 2, 2, 2, 0, 2, 2, 0, 0,  
            0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 2, 0, 0, 0, 2, 0, 0, 0, 0,  
            0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 2, 0, 0, 0,  
            0, 0, 0, 0, 0, 2, 0, 0, 0], dtype=int32)
```

```
[99]: #Obtain the centroids
cluster_center = kmeans_plants.cluster_centers_
cluster_center
```

```
[99]: array([[ 1.11418969,  1.12703719],
           [-1.27561714, -1.24966339],
           [ 0.27487366,  0.24182267]])
```

```
[100]: X_z.tail()
```

```
[100]:      petal_length_z  petal_width_z
312          1.686861      1.576880
313          0.687793      0.363339
314          1.686861      1.085716
315          0.965312      0.767852
316          0.743297      1.576880
```



```
[101]: #Merge original data, rescaled data and add column to label the results of clustering
```

```
clt = pd.DataFrame(cluster, columns=['cluster'])
plants_cluster = X.merge(X_z, on = X.index)
plants_cluster = pd.concat([plants_cluster, clt], axis = 1, sort = True)
plants_cluster = plants_cluster.rename(columns={'key_0': 'ID'})
plants_cluster.tail()
```

```
[101]:
```

	ID	petal_length	petal_width	petal_length_z	petal_width_z	cluster
312	321	6.8	2.400000	1.686861	1.576880	0
313	322	5.0	1.500000	0.687793	0.363339	2
314	323	6.8	2.035738	1.686861	1.085716	0
315	324	5.5	1.800000	0.965312	0.767852	0
316	325	5.1	2.400000	0.743297	1.576880	0

```
[102]: #Preview the new dataset
plants_cluster.head()
```


```
[102]:
```

	ID	petal_length	petal_width	petal_length_z	petal_width_z	cluster
0	0	1.9	0.299057	-1.032823	-1.255988	1
1	1	1.6	0.200000	-1.199335	-1.389553	1
2	2	1.8	0.400000	-1.088327	-1.119878	1
3	3	1.4	0.300000	-1.310342	-1.254716	1
4	4	1.6	0.200000	-1.199335	-1.389553	1

```
[103]: #Plot clusters
plt.figure(figsize=(8,6))
sns.scatterplot(plants_cluster['petal_length_z'],
                plants_cluster['petal_width_z'],
                hue=plants_cluster['cluster'])

#Plot centroids of the clusters
plt.plot(cluster_center[:,0],
         cluster_center[:,1],
         'r*',
         markersize=10)
# Centroids marked with stars
```

```
[103]: [<matplotlib.lines.Line2D at 0x12ba565f8>]
```



```
[104]: #Predict plants
#Suppose we have 5 plants and they have the following measurements
#The data is saved in a list, each element of which has two values: width and
→length
plant_list = [[4.5, 2.3],
              [2.92, 1.5],
              [1.2, 0.3],
              [1.2, 0.5],
              [4.9, 2]]
#Convert the list to a dataframe
newplant = pd.DataFrame(plant_list,columns=['petal_length','petal_width'])
newplant
```

```
[104]:   petal_length  petal_width
0         4.50         2.3
1         2.92         1.5
2         1.20         0.3
3         1.20         0.5
4         4.90         2.0
```

```
[105]: #Standardize the measurements of new plants
newplant_z = (newplant-plants_mean)/plants_std
```

```
newplant_z = newplant_z.rename(columns={'petal_length': 'petal_length_z',
    ↳ 'petal_width': 'petal_width_z'})
newplant_z
```

```
[105]:   petal_length_z  petal_width_z
0         0.409627         1.439765
1        -0.465948         0.362765
2        -1.419106        -1.252735
3        -1.419106        -0.983485
4         0.631291         1.035890
```

```
[106]: #Prediction
preds = kmeans_plants.predict(newplant_z)
preds
```

```
[106]: array([0, 2, 1, 1, 0], dtype=int32)
```

```
[107]: #Interpretations
#The first plant belongs to the cluster of high petal_length and petal_width
#The second plant belongs to the cluster of medium petal_length and petal_width
#The thrid plant belongs to the cluster of low petal_length and petal_width
#The fourth plant belongs to the cluster of low petal_length and petal_width
#The fifth plant belongs to the cluster of high petal_length and petal_width

# Combine original data, standardized data and predicted clusters
combined_newplant = pd.concat([newplant, newplant_z, pd.DataFrame(preds,
    ↳ columns=['cluster'])], axis = 1)
combined_newplant
```

```
[107]:   petal_length  petal_width  petal_length_z  petal_width_z  cluster
0          4.50          2.3         0.409627         1.439765         0
1          2.92          1.5        -0.465948         0.362765         2
2          1.20          0.3        -1.419106        -1.252735         1
3          1.20          0.5        -1.419106        -0.983485         1
4          4.90          2.0         0.631291         1.035890         0
```

```
[108]: #Plot clusters
plt.figure(figsize=(8,6))
sns.scatterplot(plants_cluster['petal_length_z'],
                plants_cluster['petal_width_z'],
                hue=plants_cluster['cluster'])


#Plot centroids of the clusters
plt.plot(cluster_center[:,0],
         cluster_center[:,1],
         'r*',
         markersize=10)

#Plot predictions
plt.plot(newplant_z.iloc[:,0],
```

```
newplant_z.iloc[:,1],
'go',
markersize=10)
```

*# Graph shows where predicted flowers should lie with green dots*

[108]: [



## 1.4 Classification

[109]: *#--Preview Data and drop non-predictors*  
plants.  
→drop(['ID', 'sepal\_length\_z', 'sepal\_width\_z', 'petal\_length\_z', 'petal\_width\_z', 'SpeciesName'])  
→axis = 1, inplace = True)  
plants.head()

[109]:

	sepal_length	sepal_width	petal_length	petal_width	species
0	4.3	2.8	1.9	0.299057	0
1	4.8	3.4	1.6	0.200000	0
2	5.0	3.0	1.8	0.400000	0
3	4.8	3.0	1.4	0.300000	0
4	5.1	3.8	1.6	0.200000	0

```
[110]: # establish predictors
outcome = 'species'
predictors = [c for c in plants.columns if c != outcome]
predictors

[110]: ['sepal_length', 'sepal_width', 'petal_length', 'petal_width']

[111]: # --Split training and testing
# Test size of 30% used which is around the industry standard
X = plants.drop('species',axis=1) # -- features --
y = plants['species']             # -- target --

x_train, x_test, y_train, y_test = train_test_split(X,y,
                                                    test_size = 0.3,
                                                    random_state=1)

[112]: # -- train neural nets and fit model--
# Solver type is lbfgs
# Hidden layer sizes is 3
ann_clf = MLPClassifier(hidden_layer_sizes = (3),
                        activation = 'logistic',
                        solver = 'lbfgs', random_state = 1)

ann_clf.fit(x_train,y_train)

[112]: MLPClassifier(activation='logistic', alpha=0.0001, batch_size='auto',
                    beta_1=0.9, beta_2=0.999, early_stopping=False, epsilon=1e-08,
                    hidden_layer_sizes=3, learning_rate='constant',
                    learning_rate_init=0.001, max_iter=200, momentum=0.9,
                    n_iter_no_change=10, nesterovs_momentum=True, power_t=0.5,
                    random_state=1, shuffle=True, solver='lbfgs', tol=0.0001,
                    validation_fraction=0.1, verbose=False, warm_start=False)

[113]: # -- confusion matrix --

metrics.confusion_matrix(y_true = y_train,
                        y_pred = ann_clf.predict(x_train))

[113]: array([[76,  0,  0],
              [ 0, 68,  1],
              [ 0,  2, 74]])

[114]: # -- Use sklearn.metrics to present confusion_matrix --

# -- use seaborn heatmap to present the confusion matrix --
# -- This is based on TRAINING DATA --
%matplotlib inline

sns.heatmap(metrics.confusion_matrix(y_true = y_train,
                                    y_pred = ann_clf.predict(x_train)),
```


```

        annot=True,
        cbar=False,
        fmt='g',
        cmap = plt.cm.get_cmap('Blues'))

plt.ylabel('Actual')
plt.xlabel('Predicted')

```

[114]: Text(0.5, 16.0, 'Predicted')



```

[115]: # -- Validation performance --
# -- Use test data --
metrics.confusion_matrix(y_true = y_test,
                          y_pred=ann_clf.predict(x_test))

```

[115]: array([[30, 0, 0],  
[ 0, 35, 2],  
[ 0, 2, 27]])

```


[116]: # --- Print Confusion Matrix Using Panda's crosstab function ---
sns.set_style('white')
print('--- Confusion Matrix (predict on rows, actual on columns ---')
cmtab = metrics.confusion_matrix(y_test,ann_clf.predict(x_test))

# confusion_matrix = pd.crosstab(preds, y_test,rownames=['Predicted'],
→colnames=['Actual'])
ax = sns.heatmap(cmtab, annot=True,fmt='g',cbar=False,cmap='Blues')
ax.set(xlabel="Predicted",ylabel="Actual");

```

```
# The confusion matrix shows that nearly all of the predicted values were
→ correctly identified. There were 2 flowers
# in classes 1 and 2 were misidentified. This is very minimal in the grand
→ scheme of the dataset
```

--- Confusion Matrix (predict on rows, actual on columns) ---



```
[117]: # show metrics
print(metrics.classification_report(y_true = y_test,
                                     y_pred = ann_clf.predict(x_test)))

# The model has high performance in all areas. Especially the accuracy,
→ precision, and F1 score which are
# good metrics for this model because our biggest concern is classifying
→ flowers correctly and avoiding false positives
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	30
1	0.95	0.95	0.95	37
2	0.93	0.93	0.93	29
accuracy			0.96	96
macro avg	0.96	0.96	0.96	96

weighted avg	0.96	0.96	0.96	96
--------------	------	------	------	----